

University of Waterloo

CS241 - Winter 2024 - Course Notes

Author: Brandon Zhou

Course Code: CS241

Course Name: Foundations of Sequential Programs

Instructor: Chengnian Sun

Section: 001

Date Created: January 09, 2024

Final Exam Date: TBA

Disclaimer: These course notes are intended to supplement primary instructional materials and facilitate learning. It's worth mentioning that some sections of these notes might have been influenced by ChatGPT, an OpenAI product. Segments sourced or influenced by ChatGPT, where present, will be clearly indicated for reference.

While I have made diligent efforts to ensure the accuracy of the content, there is a potential for errors, outdated information, or inaccuracies, especially in sections sourced from ChatGPT. I make no warranties regarding the completeness, reliability or accuracy of the notes contained in this notebook. It's crucial to view these notes as a supplementary reference and not a primary source.

Should any uncertainties or ambiguities arise from the material, I strongly advise consulting with your course instructors or the relevant course staff for comprehensive understanding. I apologize for any potential discrepancies or oversights.

Any alterations or modifications made to this notebook after its initial creation are neither endorsed nor recognized by me. For any doubts, always cross-reference with trusted academic resources.

Table of Contents

- **Section 1: Lectures**
 - [Lecture 1](#)
 - [Lecture 2](#)
 - [Lecture 3](#)
 - [Lecture 4](#)
 - [Lecture 5](#)
 - [Lecture 6](#)
 - [Lecture 7](#)
 - [Lecture 8](#)
 - [Lecture 9](#)
 - [Lecture 10](#)
 - [Lecture 11](#)
 - [Lecture 12](#)
 - [Lecture 13](#)
 - [Lecture 14](#)
 - [Lecture 15](#)
 - [Lecture 16](#)
 - [Lecture 17](#)
 - [Lecture 18](#)
 - [Lecture 19](#)
 - [Lecture 20](#)
 - [Lecture 21](#)
 - [Lecture 22](#)
 - [Lecture 23](#)
 - [Lecture 24](#)
 - **Section 2: Tutorials**
 - [Tutorial 1](#)
 - **Section 3: Reviews**
 - [Final Review](#)
-

Section 1: Lectures

Lecture 1

- Definition: A bit is a binary digit. That is, a 0 or 1 (off or on)
- Definition: A nibble is 4 bits.
 - Example: 1001
- Definition: A byte is 8 bits.
 - Example: 10011101
- in C/C++:
 - Char: 8 bits
 - Unsigned char: 8 bits
 - Short: 2 bytes/16 bits

- int: 4 bytes
- longlong: 16 bytes
- Definition: A word is a machine-specific grouping of bytes. For us, a word will be 4 bytes (32-bit architecture) though 8-byte (or 64-bit architectures) words are more common now.
- Definition (Hexadecimal Notation): The base-16 representation system is called the hexadecimal system. It consists of the numbers from 0 to 9 and the letters a, b, c, d, e, f (which convert to the numbers from 10 to 15 in decimal notation)
 - Sometimes we denote the base with a subscript like 10011101_2 and $9d_{16}$.
 - Also, for hexadecimal, you will routinely see the notation $0x9d$. (The $0x$ denotes a hexadecimal representation in computer science).
 - Note that each hexadecimal character is a nibble (4 bits).
- Conversion Table
 - Note: upper case letters are also used for hexadecimal notation. Context should make things clear.

Binary	Decimal	Hex
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7

Binary	Decimal	Hex
1000	8	8
1001	9	9
1010	10	a
1011	11	b
1100	12	c
1101	13	d
1110	14	e
1111	15	f

- Notation:
 - Binary: 0, 1
 - Decimal(Base 10): 0, ..., 9
 - Hexadecimal: 0, ..., 9, A(10), B(11), C(12), D(13), E(14), F(15)
- Examples:
 - $0000(\text{base } 2) \rightarrow 0x0(\text{base } 16)$
 - $1111(\text{base } 2) \rightarrow 0xf(\text{base } 16)$
- What do bytes represent?
 - Numbers
 - Characters
 - Garbage in memory
 - Instructions (Words, or 4 bytes, will correspond to a compute instruction in our computer system)
- Bytes as Binary Numbers
 - Unsigned (non-negative integers)
 - $b_7 \dots b_0(\text{base } 2) = b_7 \times 2^7 + \dots + b_0 \times 2^0(\text{base } 10)$
 - Example: $01010101 = 0 \times 2^7 + \dots + 1 \times 2^0$
 - Converting to Binary:

- One way: Take the largest power of 2 less than the unsigned integer, subtract and repeat
- Another way is to constantly divide by 2, get the remainder for each division, reading *from the bottom to up* at the end, and that will be the binary representation of this unsigned integer
 - Example: 38

Number	Quotient	Remainder
38	19	0
19	9	1
9	4	1
4	2	0
2	1	0
1	0	1

- Brief Explanation: Consider

$$N = b_0 + 2b_1 + 2^2b_2 + \dots$$

The remainder when dividing N by 2 gives the b_0 value. After doing $\frac{N-b_0}{2}$, we end of with

$$\frac{N-b_0}{2} = b_1 + 2b_2 + 2^2b_3 + \dots$$

and we can repeat the process. (This is why we have to read bottom-up as we get b_0 first, then $b_1 \dots$)

- Signed integers
 - Attempt 1: make the first bit a signed bit. This is called the "sign-magnitude" representation
 - Problems:
 - Two representations of 0(wasteful and awkward)
 - Arithmetic is tricky. Is the sum of a positive number and a negative number positive or negative? It depends!
 - Attempt 2: Two's complement form
 - Similar to "sign-magnitude" representation in spirit, first bit is 0 if non-negative, 1 if negative

- Negate a value by just subtracting from zero and *letting it overflow*.
- Decimal to Two's Complement:
 - A trick to the same thing of negating a value:
 - Take the complement of all bits (flip the 0 bits to 1 and 1 bits to 0)
 - Add 1
 - A slightly faster trick is to locate the rightmost 1 bit and flip all the bits to the left of it
 - Example: 11011010 Negating: 00100110 = 00100101 + 1
 - Note: Flipping the bits and adding 1 is the same as
 - subtracting 1 and flipping the bits for non-zero numbers
 - subtracting from 0
 - Example: compute -38_{10} using this notation in one byte of space:
 - Step 1: $38_{10} = 00100110_2$
 - Step 2: take the complement of all the bits: 11011001_2
 - Step 3: plus 1: 11011010_2
- Two's Complement to Decimal
 - Let's compute -38_{10} using one-byte Two's complement. First, write 38 in binary: $38_{10} = 00100110_2$. Next, take the complement of all the bits 11011001_2 . Finally, add 1: 11011010_2 . This last value is -38_{10} .
 - To convert 11011010_2 , a number in Two's complement representation, to decimal, one method is to flip the bits and add 1: $00100110_2 = 2^5 + 2^2 + 2^1 = 38$. Thus, the corresponding positive number is 38 and so the original number is -38 .
 - Another way to do this computation is to treat the original number 11011010_2 as an unsigned number, convert to decimal and subtract 256 from it (since we have 8 bits, and the first bit is a 1 meaning it should be a negative value). This also gives -38

$$\begin{aligned}
 11011010_2 &= 2^7 + 2^6 + 2^4 + 2^3 + 2^1 - 2^8 \\
 &= 128 + 64 + 16 + 8 + 2 - 256 \\
 &= 218 - 256 \\
 &= -38
 \end{aligned}$$

- The idea behind [one byte] Two's Complement notation is based on the following observations:
 - The range for unsigned integers is 0 to 255. Recall that 255 is 11111111_2 . If we add 1 to 255, then, after discarding overflow bits, we get the number 0.
 - Thus, let's treat 2^8 as 0, i.e., let's work modulo $2^8 = 256$. In this vein, we set up a correspondence between the positive integer k and the unsigned integer $2^8 - k$. Since we are working modulo 2^8 , subtracting a positive integer k from 0 is the same as subtracting it from 2^8 .
 - In this case, note that $255 = 2^8 - 1 = 2^7 + 2^6 + 2^5 + 2^4 + 2^3 + 2^2 + 2^1 + 2^0$ and in general

$$2^n - 1 = \sum_{i=0}^{n-1} 2^i$$

- As an explicit example (which can be generalized naturally) take a number, say $38_{10} = 00100110_2 = 2^5 + 2^2 + 2^1$. What should the corresponding negative number be? Well, note that we've said subtracting a positive integer k from 0 is the same as subtracting it from 2^8 :

$$\begin{aligned}
 2^8 - 1 &= 2^7 + 2^6 + 2^5 + 2^4 + 2^3 + 2^2 + 2^1 + 2^0 \\
 2^8 - 1 &= 38 + 2^7 + 2^6 + 2^4 + 2^3 + 2^0 \\
 2^8 - 38 &= 2^7 + 2^6 + 2^4 + 2^3 + 2^0 + 1 \quad \text{(flip the bits and add 1)}
 \end{aligned}$$

- We mentioned that another method of negating a two's complement number is to flip the bits to the left of the rightmost 1. Justify why this technique works.
 - Every bit to the right of the rightmost 1 is a 0. When we "flip the bits", these become 1s. When we "add 1", the carry propagates up until the position of the "rightmost 1" (the "rightmost 1" is 0 after flip and will stop propagating when the carry reaches this point, and everything on the left of the "rightmost 1" is flipped).
- The main difference between signed and unsigned binary arithmetic is that we are now working modulo 256 (or, more generally, 2^n in the case of n-bit Two's complement numbers)
- When working in Two's complement, overflow occurs when adding numbers if the original two numbers have the same sign, but the result has a different sign.
- Arithmetic of Signed Integers
 - All of the arithmetic works by ignoring overflow precisely because arithmetic works in \mathbb{Z}_{256} !
- What is the range of numbers expressible in one-byte Two's Complement notation?
 - $-128 \sim 127$
- Definitions: The Most Significant Bit (MSB) is the left- most bit (highest value/sign bit); The Least Significant Bit (LSB) is the right-most bit (lowest value).

Lecture 2

- ASCII (American Standard Code for Information Interchange) uses 7 bits to represent characters.

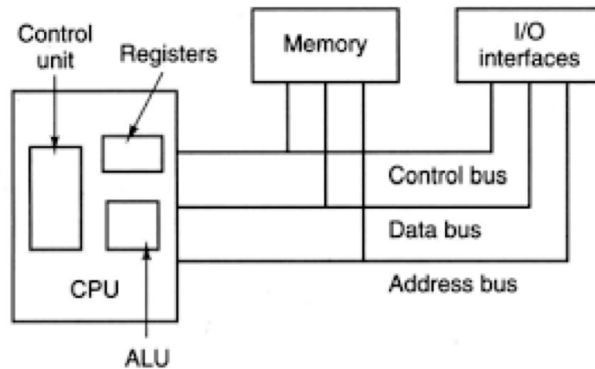
```

char c = '0';
printf("%c", c);
// stdout: 0
printf("%d", (int)c);
// stdout: 48

```

- Bit-Wise Operators

- Example: suppose we have `unsigned char a=5, b=3;`, which means $a = 00000101, b = 00000011$
 - Bitwise not $\sim a$, for example $c = \sim a$; gives $c = 11111010$
 - Bitwise and $\&$, for example $c = a \& b$; gives $c = 00000001$
 - Bitwise or $|$, for example $c = a | b$; gives $c = 00000111$
 - Bitwise exclusive \wedge (only true if the bit in a and b is different), for example $c = a \wedge b$; gives $c = 00000110$
 - Bitwise shift right or left $>>$ and $<<$, for example
 - $c = a >> 2$; gives $c = 00000001$ and
 - $c = a << 3$; gives $c = 00101000$
 - $a << 1$ equivalent to $a * 2$
 - $a >> 1$ equivalent to $a / 2$
 - These can even be combined with the assignment operator: $c \&= 5;$



- CPU with Memory
- MIPS has 32 registers that are called “general purpose”
 - Some general-purpose registers are special:
 - $\$0$ is always 0
 - $\$31$ is for return address
 - $\$30$ is our stack pointer
 - $\$29$ is our frame pointer
- Problem: We only know from context what bits have what meaning, and in particular, which are instructions.
 - Solution: Convention is to set memory address 0 in RAM to be an instruction. $0x0$: instruction $i1$
- Problem: How does MIPS know what to do next?
 - Solution: Have a special register called the Program Counter (or PC for short) to tell us what instruction to do next.
- Problem: How do we put our program into RAM?
 - Solution: A program called a loader puts our program into memory and sets the PC to be the first address.
- Algorithm 1 Fetch-Execute Cycle

```
PC=0
while true do
  IR = MEM[PC]
  PC += 4
  Decode and execute instruction in IR
end while
```

- Write a program in MIPS that adds the values of registers $\$8$ and $\$9$ and stores the result in register $\$3$.

```
add $d, $s, $t
0000 00ss ssst tttt dddd d000 0010 0000
```

- Why 5 bits for each? Because there are 32 registers $\$0 - \31 and $2^5 = 32$
- Adds registers $\$s$ and $\$t$ and stores the sum in register $\$d$. Important! The order of $\$d, \s and $\$t$ are shifted in the encoding.

Lecture 3

- **Putting values in registers** Load immediate and skip. This places the next value in RAM [an immediate] into $\$d$ and increments the program counter by 4 (it skips the next line which is usually not an instruction).

```
lis $d: 0000 0000 0000 0000 dddd d000 0001 0100
# What it really does is
$d = MEM[PC]
PC = PC + 4
```

- How do we get the value we care about into the next location in RAM?

```
.word i: iiii iiii iiii iiii iiii iiii iiii iiii
```

- The above is an assembler directive (not a MIPS instruction). The value i , as a two's complement integer, is placed in the correct memory location in RAM as it occurs in the code.
 - Can also use hexadecimal values: $0xi$
 - Decimal is also allowed.
- Example:

```
lis $1
.word 10
# At this moment, $1 = 10
add $1, $1, $0
```

- Example: Write a MIPS program that adds together 11 and 13 and stores the result in register \$3.

```
lis $1
.word 11
lis $2
.word 13
add $3, $1, $2
# Solution on the course notes
lis $8      0000 0000 0000 0000 0100 0000 0001 0100
.word 11    0000 0000 0000 0000 0000 0000 0000 1011
lis $9      0000 0000 0000 0000 0100 1000 0001 0100
.word 0xd   0000 0000 0000 0000 0000 0000 0000 1101
add $3,$8,$9 0000 0001 0000 1001 0001 1000 0010 0000
# The code on the left is what we call Assembly Code.
# The code on the right is what we call Machine Code.
```

- Jump Register. Sets the pc to be \$s.

```
jr $s
0000 00ss sss0 0000 0000 0000 0000 1000
```

- For us, our return address will typically be in \$31, so we will typically call the below. This command returns control to the loader.

```
jr $31
0000 0011 1110 0000 0000 0000 0000 1000
```

- So the complete example for the example is (so that the while loop will terminate)

```
lis $8      0000 0000 0000 0000 0100 0000 0001 0100
.word 11    0000 0000 0000 0000 0000 0000 0000 1011
lis $9      0000 0000 0000 0000 0100 1000 0001 0100
.word 0xd   0000 0000 0000 0000 0000 0000 0000 1101
add $3,$8,$9 0000 0001 0000 1001 0001 1000 0010 0000
jr $31      0000 0011 1110 0000 0000 0000 0000 1000
```

- To multiply two words, we need to use the two special registers `hi` and `lo`.
 - `hi` is most significant 4 bytes
 - `lo` is least significant 4 bytes

```
mult $s, $t
0000 00ss ssst tttt 0000 0000 0001 1000
```

- The above performs the multiplication and places the most significant word (largest 4 bytes) in `hi` and the least significant word in `lo`.
- `div $s, $t` performs integer division and places the quotient s/t in `lo` [`lo quo`] and the remainder `st` in `hi`. Note the sign of the remainder matches the sign of the divisor stored in `$s`.

```
div $s, $t
0000 00ss ssst tttt 0000 0000 0001 1010
```

- Multiplication and division happen on these special registers `hi` and `lo`. How can I access the data?
 - Move from register `hi` into register `$d`.

```
mfhi $d
0000 0000 0000 0000 dddd d000 0001 0000
```

- Move from register `lo` into register `$d`.

```
mflo $d
0000 0000 0000 0000 dddd d000 0001 0010
```

- RAM
 - Large[r] amount of memory stored off the CPU.
 - RAM access is slower than register access (but is larger, as a tradeoff).
 - Data travels between RAM and the CPU via the bus.
 - Modern day RAM consists of in the neighbourhood of 10^{10} bytes.
 - Instructions occur in RAM starting with address 0 and increase by the word size (in our case 4).
 - But, this simplification will vanish later...
 - Each memory block in RAM has an address; say from 0 to $n - 1$
 - Words occur every 4 bytes, starting with byte 0. Indexed by 0, 4, 8, ... $n - 4$.
 - Words are formed from consecutive, aligned (usually) bytes.
 - Cannot directly use the data in the RAM. Must transfer first to registers.
- Load word. Takes a word from RAM and places it into a register. Specifically, load the word in `MEM[$s + i]` and store in `$t`.

```
lw $t, i($s)
1000 11ss ssst tttt iiii iiii iiii iiii
# which is equivalent to
$t = MEM[$s + i]
# Example
lw $1, -4($30)
```

```
# which means
$t1 <- MEM[$$0 - 4]
```

- **Store word.** Takes a word from a register and stores it into RAM. Specifically, load the word in $\$t$ and store it in $\text{MEM}[\$s + i]$.

```
sw $t, i($s)
1010 11ss ssst tttt iiii iiii iiii iiii
```

- Note that i must be an immediate, NOT another register! It is a 16-bit Two's complement immediate
- Example: Suppose that $\$1$ contains the address of an array of words, and $\$2$ takes the number of elements in this array (assume less than 220). Place the number **7** in the last possible spot in the array.

```
# First element in the array is arr, then the second element is arr + 4 ... the last element in the array is arr + 4 * (length - 1)
lis $8 ; 7
.word 7
lis $9 ; 4
.word 4
mult $2, $9 ; length * 4
mflo $3 ; length * 4
add $3, $3, $1 ; arr + length * 4
sw $8, -4($3) ; MEM[$3 - 4] = $8
jr $31
```

- **Branch on equal.** If $\$s == \t then $\text{pc} += i * 4$. That is, skip ahead i many instructions if $\$s$ and $\$t$ are equal.

```
beq $s, $t, i
0001 00ss ssst tttt iiii iiii iiii iiii
# It is like
if ($s == $t) {
    PC += i*4
}
```

- **Branch on not equal.** If $\$s != \t then $\text{pc} += i * 4$. That is, skip ahead i many instructions if $\$s$ and $\$t$ are not equal.

```
bne $s, $t, i
0001 01ss ssst tttt iiii iiii iiii iiii
# It is like
if ($t != $s) {
    PC += i*4
}
```

- Example:

```
beq $0, $0, 0 ; This executes the next instruction as PC has been updated to +4 already
beq $0, $0, 1 ; This executes the second next instruction
```

Lecture 4

- Write an assembly language MIPS program that places the value **3** in register $\$2$ if the signed number in register $\$1$ is odd and places the value **11** in register $\$2$ if the number is even.

```
lis $8 ; $8 = 2
.word 2
lis $9 ; $9 = 3
.word 3
lis $2 ; $2 = 11
.word 11
div $1, $8
mfhi $3
beq $3, $0, 1
add $2, $9, $0
jr $31
```

- **Set Less Than.** Sets the value of register $\$d$ to be **1** provided the value in register $\$s$ is less than the value in register $\$t$ and sets it to be **0** otherwise.

```
slt $d, $s, $t
0000 00ss ssst tttt dddd d000 0010 1010
# which basically means
if ($s < $t) {
    $d = 1 (true)
} else {
    $d = 0 (false)
}
```

- Example: Write an assembly language MIPS program that negates the value in register $\$1$ provided it is positive.

```
slt $2, $1, $0
bne $2, $0, 1
sub $1, $0, $1
jr $31
```

- Exercise: Write an assembly language MIPS program that places the absolute value of register **1** in register **2**.

```
add $2, $1, $0 ; $2 = $1
slt $3, $0, $1 ; 0 < $1
bne $3, $0, 1
```

```
sub $2, $0, $2 ; $2 = 0 - $2
jr $31
```

- Looping exmaple: Write an assembly language MIPS program that adds together all even numbers from 1 to 20 inclusive. Store the answer in register \$3.
 - Note: semicolons for comments in MIPS assembly

```
lis $2
.word 20
lis $1
.word 2
add $3, $0, $0
add $3, $3, $2 ; line -3
sub $2, $2, $1 ; line -2
bne $2, $0, -3 ; line -1 from here
jr $31
```

- Labels aren't machine code, so don't take words. That means that for `beq` and `bne`, *labels don't have "line numbers" on their own*. A label at the end of code is allowed. It has the address of what would be the first instruction after the program.

```
label: operation commands
```

- Example: `sample` has the address `0x4`, which is the location of `add $1, $0, $0`.

```
sub $3, $0, $0
sample:
add $1, $0, $0
```


- A better way to loop without hard-coding `-3` in the previous example is (otherwise, if we were to, say, add a new instruction in between the lines specified by our branching, all our numbers would be incorrect.)
 - Note that `top` in `bne` is computed by the assembler to be the *difference between the program counter and top*. That is, here it computes $(top - PC)/4$ which is $(0x14 - 0x20)/4 = -3$
 - `PC` is the line number after the current line

```
lis $2
.word 20
lis $1
.word 2
add $3, $0, $0
top:
    add $3, $3, $2
    sub $2, $2, $1
    bne $2, $0, top
jr $31
```


- RAM



- Register `$30` initially points to the very bottom of the free RAM. It can be used as a bookmark to separate the used and unused free RAM if we allocate from the one end, and push and pop things like a stack! In other words, we will use `$30` as a pointer to the top of a stack.
- Really, `$30` points to the top of the stack of memory in RAM.




- Because our program is at zero, the stack grows from high memory to low memory, so pushing involves reducing the value of `$30`.



- Example: Suppose procedures f , g and h are such that:

```
f calls procedure g
  g calls procedure h
    h returns
  g returns
f returns.
```



- In the previous example:
 - Calling procedures pushes more registers onto the stack and returning pops them off.
 - This is a stack, and we call \$30 our **stack pointer**.
 - We can also use the stack for local storage if needed in procedures. Just reset \$30 before procedures return.
- Template for a procedure f that modifies registers \$1 and \$2:

```
f:
sw $1, -4($30) ; Push registers we modify
sw $2, -8($30)
lis $2 ; Decrement stack pointer
.word 8
sub $30, $30, $2
; Insert procedure here
add $30, $30, $2 ; Assuming $2 is still 8
lw $2, -8($30) ; Pop = restore
lw $1, -4($30)
; Uh oh! How do we return?
```

- There is a problem with returning:

```
main:
lis $8
.word f ; Recall f is an address
jr $8 ; Jump to the first line of f
(NEXT LINE)
```

- Once f completes, we really want to jump back to the line labelled above as (NEXT LINE), i.e., set the program counter back to that line. How do we do that?
- Jump and Link Register. Sets \$31 to be the PC and then sets the PC to be \$s. Accomplished by `temp = $s then $31 = PC then PC = temp`.

```
jblr $s
0000 00ss sss0 0000 0000 0000 1001
```

- Main Changes

```
main:
lis $8
.word f
sw $31, -4($30) ; Push $31 to stack
lis $31 ; Use $31 since it's been saved .word 4
sub $30, $30, $31
jblr $8 ; Overwrites $31
lis $31 ; Use $31 since we'll restore it .word 4
```



```

add $30, $30, $31
lw $31, -4($30) ; Pop $31 from stack
jr $31 ; Return to loader

```

- Procedure Changes

```

f:
    sw $1, -4($30) ; Push registers we will modify
    sw $2, -8($30)
    lis $2
    .word 8
    sub $30, $30, $2 ; Decrement stack pointer
    ; Insert procedure here
    add $30, $30, $2 ; Assuming $2 is still 8
    lw $2, -8($30) ; Pop registers to restore
    lw $1, -4($30)
    jr $31 ; New line!

```

Lecture 5

- Note: there is NO default value to register, so remember to initialize it
- How do we pass parameters?

```
void f(int a, int b) {}
```

- Example: sumEvens1ToN adds all even numbers from 1 to N
 - \$1 Scratch Register (Should Save!)
 - \$2 Input Register (Should Save!)
 - \$3 Output Register (Do NOT Save!)

```

; The idea is:
; $3 = 0
; $1 = $2 % 2
; $2 = $2 - $1
; $1 = 2
; Top: $3 = $3 + $2
;     $2 = $2 - 2
;     if ($2 != 0) go to top
; jr $31

```

```

lis $1
.word 8
sw $1, -4($30)
sw $2, -8($30)
sub $30, $30, $1
add $3, $0, $0
lis $1
.word 2
div $2, $1
mfhi $1
sub $2, $2, $1
lis $1
.word 2
top:
add $3, $3, $2
sub $2, $2, $1
bne $2, $0, top
lis $1
.word 8
add $30, $30, $1
lw $2, -8($30)
lw $1, -4($30) ; Reload $1 and $2
jr $31 ; Back to caller

```

- Input and Output
 - We do this one byte at a time!
 - Output: Use `sw` to store words in location `0xffff000c`. Least significant byte will be printed.
 - Input: Use `lw` to load words in location `0xffff0004`. Least significant byte will be the next character from `stdin`.
 - Input/Output the ASCII character

```

lis $1
.word 0xffff000c
lis $2
.word 48 ; In ASCII code, 48 means 0!
sw $2, 0($1)

```

- Example: Printing CS241 to the screen followed by a newline character:

```

lis $1
.word 0xffff000c
lis $2
.word 67 ; C
sw $2, 0($1)
lis $2
.word 83 ; S
sw $2, 0($1)
lis $2
.word 50 ; 2
sw $2, 0($1)
lis $2
sw $2, 0($1)
lis $2

```

```
.word 49 ; 1
sw $2, 0($1)
lis $2
.word 10 ; \n
sw $2, 0($1)
jr $31
```

- Part of our long-term goal is to convert assembly code (our MIPS language) into machine code (bits).
 - Input: Assembly code
 - Output: Machine code
- Any such translation process involves two phases: Analysis and Synthesis.
 - Analysis: Understand what is meant by the input source
 - Synthesis: Output the equivalent target code in the new format
- What if a label is used before it is defined? We don't know the address when it's used!
 - Perform two passes:
 - Pass 1: Group tokens into instructions and record addresses of labels (data structure?).
 - Note: multiple labels are possible for the same line! For example, `f: g: add $1, $1, $1.`
 - Pass 2: translate each instructions into machine code. If it refers to a label, look up the associated address compute the value.
- A label at the end of code is allowed (it would be the address of the first line after your program).
- Our instruction (`bne`) can be broken down as follows:

Opcode
(6 bits)

Registers
(5 bits)

Regis
(5 b

- We can use bit shifting to put information into the correct position, and use a bitwise or to join them:

```
int instr = (5 << 26) | (2 << 21) | (0 << 16) | offset
```

- Recall in C++, ints are 4 bytes. We only want the last two bytes. First, we need to apply a “mask” to only get the last 16 bits:

```
offset = -3 & 0xffff
```

- Printing Bytes in C++

```
int instr = (5 << 26) | (2 << 21) | (0 << 16) | (-3 & 0xffff); unsigned char c = instr >> 24;
cout << c;
c = instr >> 16;
cout << c;
c = instr >> 8;
cout << c;
c = instr;
cout << c; // will output the least significant 8 bits
```

- Note: You can also mask here to get the 'last byte' by doing `& 0xff` if you're worried about which byte will get copied over.

Section 2: Tutorials

Tutorial 1

- What is Binary?
 - Binary - ways our machines encode info, and $b \in \{0, 1\}$
- Eg. what is 1000 be
 - $2^3 = 8$ unsigned magnitude
 - -8 2's complement
 - $[T, F, F, F]$ array of booleans
 - "backspace char" in ASCII
 - Representation matters**
- Unsigned binary: n-bit binary number is represented as $b_{n-1}, b_{n-2}, \dots, b_0$, where $b \in \{0, 1\}$
- To convert to decimal: $2^{n-1} \times b_{n-1} + \dots + 2^0 \times b_0$
- Decimal to binary
 - Idea 1: take the highest powers of 2 from the decimal (inefficient)
 - Idea 2: Repeatedly divide the number by 2, tracking the quotient & remainder
 - Eg. Convert 23 to binary

Number	Quotient	Remainder
23/2	11	1
11/2	5	1

Number	Quotient	Remainder
5/2	2	1
2/2	1	0
1/2	0	1

- and read from bottom to up, that will be $10111_2 = 23_{10}$

- 2's complement
 - Range of values for n-ary
 - Unsigned Binary $0 \sim 2^n - 1$
 - 2's complement $-2^{n-1} \sim 2^{n-1} - 1$
- Convert decimal to 2's complement
 - if number is ≥ 0 : use the unsigned representation
 - if number is < 0 :
 - Get the binary rep of the positive number
 - flip the bits
 - then add 1
- Convert from 2's complement to decimal
 - Method 1: if $b_{n-1} = 1$, then flip the bits, add 1 and negate the positive decimal
 - Method 2: Treat $b_{n-1} = 1$ as -2^{n-1} , and add the rest as unsigned representation
- Assembly
 - CS 241: MIPS
 - Runs programs \$ stores its data all in MEM (RAM)
 - 32 bits system where instructions are encoded as 4 bytes (1 word)
 - Registers hold 1 word of info
 - Special registers
 - 0 = 0\$, immutable
 - 31\$, end address in RAM, jr \$31 means return address
 - 3, 29, 30\$
 - *iii* ... <- 2's complement number
 - divu, multu, addi ... treats the register values as unsigned binary
 - **Programs live in the same space in MEM (RAM) as as the data they operate on**
 - PC cannot distinguish the two
 - Fetch-Execute Cycle

```
PC = 0x00
while True do:    // until PC = $31
  IR = MEM[PC]
  PC = PC + 4
  ... decode $ execute IR ...
done
```

- Constant Values
 - Use the Load Immediate Skip command (`lis $s`) followed by an instruction to save into `$s`.
 - `lis $s` -> skip the next instruction -> store that instruction into `$s`
 - Use with `.word i` to store `i` into `$s`
 - Eg. Store 10 into `$5`

```
lis $5
.word 10
; the above two lines of commands skips the .word 10 & sets $5 = 10
```

- Machine Code (based on the order of machine code horizontally for `lis $5`)
 - 000000 (operating code)
 - 00000 (`$s`)
 - 00000 (`$t`)
 - 00101 (`$d`)
 - 00000 (dead code, always be 0)
 - 010100 (function code)
- Machine Code (`.word 10`)
 - 00000 000 001010

Section 3: Reviews